

PA171 - Project documentation

Michal Jankovič

December 5, 2021

This document describes a (lossy and lossless) compression scheme designed and implemented as a project for PA171 Digital Image Filtering.

1 Compression scheme

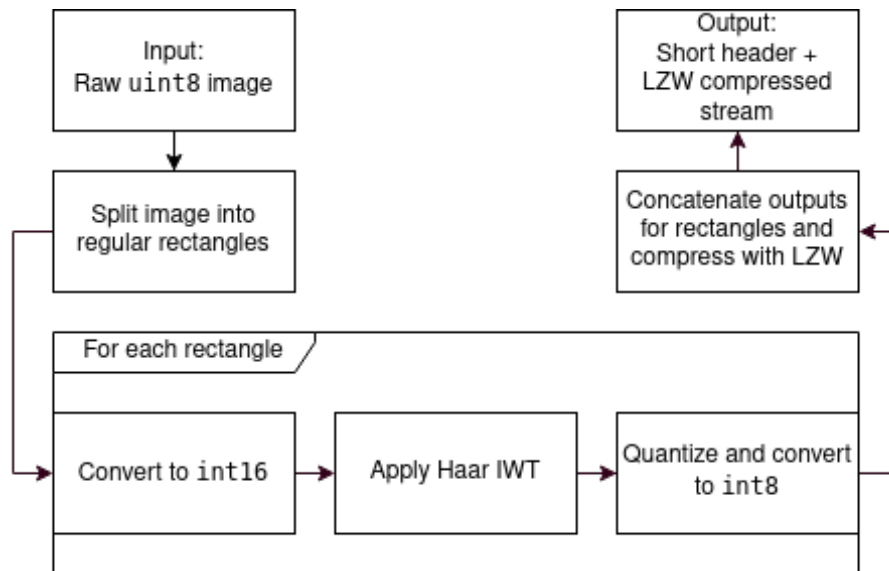


Figure 1: Compression scheme

The lossy compression scheme is based around the Haar integer wavelet transform. Input `uint8` image with values in the $[0, 255]$ range is split into regular rectangular chunks (chunks touching the right and bottom border have different sizes). These are then converted to `int16` and passed through the wavelet transform recursively until only one approximation coefficient remains.

Resulting wavelet coefficients are transformed to fit inside the single-byte `int8` range of $[-128, 127]$, and to enable better compression ratios. The 2D wavelet

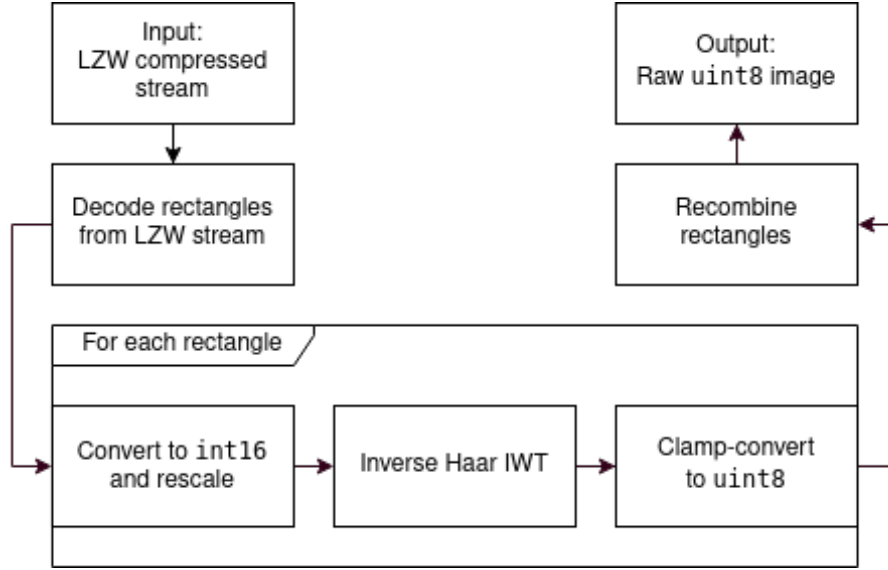


Figure 2: Decompression scheme

transform (whether it is applied just once or recursively) yields 3 kinds of coefficients that must be handled differently:

- Approximation coefficients - obtained by averaging neighboring pixels and thus in the same range as source values - $[0, 255]$. They are not quantized as they contain the most important information. Only transformation applied is a shift by -128 to fit in the `int8` range.
- Vertical / horizontal detail coefficients - obtained by differentiating neighboring pixels, these have double the range of source values - $[-256, 255]$. To fit the `int8` range, they must be quantized by dividing by at least a factor of 2. A larger factor is chosen for the detail components obtained by the first few iterations of the wavelet transform.
- Diagonal detail coefficients - obtained by twice differentiating neighboring pixels, these have four times the range of source values - $[-512, 511]$. They are quantized by twice the quantization factor for the vertical / horizontal coefficients.

Outputs from transform and quantization of each rectangle are concatenated and compressed with LZW, using all possible 8-bit values as the initial dictionary. The LZW dictionary grows up-to a maximum word size of 16 bits. When the maximum word size is reached, the dictionary is flushed.

A short header (88 bytes) is attached to the compressed LZW stream, indicating

image dimensions and all the compression parameters to allow configuring the decoder for decompression.

2 CLI application

The compression / decompression is implemented as a CLI application.

Compressing an image with default lossy compression:

```
pa171.compress image.bmp image.small
```

Compressing an image with lossless compression:

```
pa171.compress -l0 image.bmp image.small
```

Compressing an image with lossy compression level 32:

```
pa171.compress -l32 image.bmp image.small
```

Compression levels range from 0 to 64. 0 means that the input is only passed through LZW compression without transform and quantization (lossless compression). Otherwise, a quantization factor of 2x the compression level is used for the horizontal / vertical details of the first level of wavelet decomposition and 4x for diagonal details. This factor is divided by 8 for the detail components of the second level, and so on until a minimum quantization factor of 2 for vertical / horizontal details and 4 for diagonal.

Decompressing an image (with any compression level):

```
pa171.decompress image.small image.bmp
```

2.1 Compilation

Requirements:

- gcc >= 10.2.0
- CMake >= 3.17
- conan >= 1.43.0 (pip install conan)

```
mkdir build && cd build
conan install .. -b missing
cmake .. -DCMAKE_TOOLCHAIN_FILE=conan_toolchain.cmake
cmake --build .
```

Resulting binaries are in `build/bin`.